


Enterprise engineering studio · Intelligent software for modern commerce

Intelligent systems built to quietly run commerce.

Sanad constructs operational backbones and bespoke autonomous logic for commercial enterprises. We remove the cognitive and administrative weight of scaling — engineering a structural reduction of at least 30% in daily operational load.

Request Technical Dossier →

Read the 30% baseline



Enterprise data streams gather in the Sanad Core, then synchronize out to infrastructure the client controls.

For Leadership

Predictable velocity, margin protection, and operational clarity.

For Engineering


Decoupled architecture, local-first integrity, and zero black-box dependencies.

01 Proof by metric

The 30% efficiency baseline

Operational load is not a feeling. It is waiting, re-keying and reconciling — three things that can be measured, and removed. Every system we build is specified against one baseline: a structural reduction of at least 30% in daily operational load.

Workstream	Before	With Sanad
01 Inventory & Supply Sync	18 hours of multi-channel lag between stock and sales	Sub-second synchronization across every channel
02 Financial & Reconciled Auditing	Repeated reconciliations, performed by staff each cycle	Zero-loss balancing routines, automated end to end
03 Human Error Overhead	Re-keyed data entered again in each system it touches	Eliminated structurally: one entry, one record of truth



02 Systems & architecture

Ready-made modules. Bespoke engineering.

Two explicit offerings, never blended. Pillar A is production-tested software, configured for your business. Pillar B is engineered from scratch, for operational flows that belong to you alone.

A

Pre-configured SaaS · Pillar A

Turnkey Intelligent Modules

High-efficiency, production-tested systems, ready for deployment.

Select a module. Follow the data.

InventoryUnitsLedger

Warehouse scanners

3PL & logistics

Sanad Core

Client-controlled infrastructure

Owns

Stock position per node, reservations and transfers.

One stock position across every warehouse, store and channel.

Schema

```
type StockLevel {
  sku      string
  node     NodeId // warehouse or store
  on_hand  int
  reserved int    // held for open orders
  version  Lamport // deterministic ordering
}
```

Parameters

Topology	n-node, multi-warehouse
Sync	event-sourced, deterministic merge
Reservations	per-channel holds, auto-expiry
Offline	node-local ledger, reconcile on reconnect
Latency	< 100 ms in-site, sub-second across channels

A.1 · Interface wireframe

Inventory & Distributed Warehousing Logic

Stock	Transfers	Reservations	Transfer	
SKU	WH-A	WH-B	WH-C	Total
SKU-1042	128	64	0	192
SKU-2210	12	40	33	85
SKU-3087	220	0	18	238
SKU-4410	7	7	7	21
SKU-5002	54	90	31	175

● Synced 0.4 s ago5 SKUs · 3 nodes

A.2 · Interface wireframe

Hospitality & Unit Management Engines

Units	Housekeeping	This week
101	MonTueWedThuFriSatSun	
102		
103		
201		
202		

● Occupied● Turnover○ Blockedno overlaps

A.3 · Interface wireframe

Unified Commerce & POS Ledger Sync

Sync	Close	Live	
POS #A-10441	Account	Dr	Cr
	1100 Cash	46.78	-
	4200 Sales	-	42.50
VAT	4.25	2200 VAT	4.25
Total	46.79	2 Dr × 2 Cr	46.79




● Balanced, posted onceIdempotent

B

Bespoke systems · Pillar B

Tailored Commercial Architecture

From-scratch systems, engineered for unique, proprietary operational flows.

B.1  Custom internal operating systems Replacing patchwork spreadsheets One system of record in place of a folder of spreadsheets: role-based views, enforced workflows, and a history of every change. record single source of truth access role-based history append-only	B.2  Private agentic workflows Operating securely on client infrastructure Autonomous logic that runs inside your perimeter, on your hardware or your cloud account. Every action is logged, bounded by explicit permissions, and reversible. runtime client infrastructure egress denied by default oversight approval gates, full action log	B.3  High-availability distributed databases With offline-first synchronization Replicated stores that keep accepting writes when the network does not, then converge deterministically when it returns. writes local-first merge deterministic single point of failure none, by design
---	--	---

03 System mechanics & technical rigor

Specified like infrastructure, stated so it can be tested.


For CTOs and technical evaluators

Data ownership Sovereign Complete sovereignty. Databases synchronize with client-controlled infrastructure. Latency < 100 ms Sub-100ms internal state changes across multi-terminal setups.	Reliability Offline-capable Operational nodes designed to function without internet disruption. Integrations Deterministic Clean, deterministic APIs connecting legacy ERPs, payment gateways, and third-party logistics.
--	---

04 Security & compliance

Stated plainly, so it can be checked.

Properties of the systems we build, written as specifications rather than assurances.

 We do not claim certifications on this page. If you must meet a specific regulation, we design to it and show you how.	✓Data location Client-controlled infrastructure. ✓Access Role-based. Permissions are explicit and reviewable. ✓History An append-only record of changes and actions. ✓Agents Bounded permissions, approval gates, a full action log, and every action reversible. ✓Network Outbound access denied by default. ✓Failure Offline-capable nodes and deterministic recovery.
--	--

05 How we work

Four stages. One owner at each.

Every engagement follows the same order, so you always know what is being decided and what comes next.

01 Understand We map how work actually moves through your business: systems, hand-offs, waiting and re-entry. You receive A written map of current operations, and the baseline for the 30% target.	02 Design We specify the system: data model, integrations, failure behavior, and what must run offline. You receive An architecture document you can hand to your own engineers.	03 Build We build and test in your environment, in stages, each one usable on its own. You receive Working software, tested against the baseline.	04 Operate We run it with you, monitor it, and hand over ownership when you are ready. You receive Runbooks, access, and the option to run it yourselves.
---	--	---	---

06 Engagement models

Two ways to start.

	Pillar A Turnkey modules	Pillar B Bespoke systems
Best when	Your process matches a standard shape: stock, units, a sales ledger.	Your operations are particular to you and no product fits.
What we do	Configure, connect and deploy a tested module.	Specify and engineer from scratch.
From you	Access to current systems, a named process owner, and your configuration decisions.	Time with the people who run the work, access to systems and data, and a named decision-maker.
You get	A running system, and training.	A system built to your flow, with full documentation.
Start with	The technical dossier.	An architectural walkthrough.

07 Questions from engineers

Plain answers to the usual questions.

Longer answers, with diagrams, are in the technical dossier.	Who owns the data? You do. Databases synchronize with infrastructure you control, and we keep no copy that we depend on.
Ask us directly →	What happens without internet? +
	How do we migrate from our current system? +
	Can it connect to our ERP and payment gateway? +
	Do agents send our data to outside services? +
	What if we outgrow a turnkey module? +

08

The Sanad manifesto



Software should clarify the mind, not demand it.

Most enterprise software creates its own bureaucracy. Teams spend hours managing the tool that was supposed to save them time. Sanad is built on the inverse thesis: intelligent software works silently in the background so leaders and builders can dedicate their focus to what actually matters.

Support for the people who lead.

If the work is heavier than it should be, we should talk.

Request Technical Dossier →

or write to inquiries@sanad.systems

Support for the people who lead.