Enterprise engineering studio · Intelligent software for modern commerce

Intelligent systems built to quietly run commerce.



Sanad constructs operational backbones and bespoke autonomous logic for commercial enterprises. We remove the cognitive and administrative weight of scaling — engineering a structural reduction of at least 30% in daily operational load.

Request Technical Dossier → Read the 30% baseline

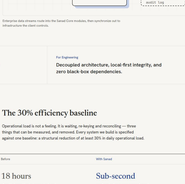
≥ 30%
Structural reduction in daily operational load

18 h → < 1 s
Inventory and supply synchronization

< 100 ms
Internal state-change across terminals

Zero
Black-box dependencies

FIG. 1 Data routing into Sanad Core



Enterprise data streams route into the Sanad Core modules, then synchronize out to infrastructure the client controls.

For Leadership
Predictable velocity, margin protection, and operational clarity.

For Engineering
Decoupled architecture, local-first integrity, and zero black box dependencies.

01 Proof by metric

The 30% efficiency baseline

Operational load is not a feeling. It is waiting, re-keying and reconciling — three things that can be measured, and removed. Every system we build is specified against one baseline: a structural reduction of at least 30% in daily operational load.

Workstream	Before	With Sanad
01 Inventory & Supply Sync	18 hours of multi-channel lag between stock and rates	Sub-second synchronization across every channel.
02 Financial & Reconciled Auditing	Repeated reconciliation, performed by staff each cycle	Zero-loss balancing routines, automated end to end
03 Human Error Overhead	Re-keyed data entered again in each system it touches	Eliminated structurally: one entry, one record of truth

Definition

```
L = T_wait + T_rekey + T_reconcile
daily operational load, in hours

C = 0.70 × L
the 30% Law
```

Daily operational load, indexed — lower Line = 100



02 Systems & architecture

Ready-made modules. Bespoke engineering.

Two explicit offerings, never blended. Pillar A is production-tested software, configured for your business. Pillar B is engineered from scratch, for operational flows that belong to you alone.

A Pre-configured Tool · Pillar A
Turnkey Intelligent Modules

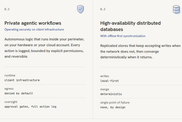
High-efficiency, production-tested systems, ready for deployment.

Select a module. Follow the data.

Choose Inventory, Units or Ledger to trace which sources feed it, which records it owns, and what it sends to your infrastructure.

- **A.1 Inventory & Distributed Warehousing Logic** — The brain platform with live rates, distributed, clone and channel.
- **A.2 Hospitality & Unit Management Engines** — Units, stays and turnover, held to a single source of availability.
- **A.3 Unified Commerce & POS Ledger Sync** — Every sale, online or in store, posts once to one ledger.

FIG. 2 Route Inventory & Distributed Warehousing Logic



Data
Stock position per node, reservations and transfers.

Parameters

Topology	n-node, multi-warehouse
Sync	event-sourced, deterministic merge
Reservations	per-channel holds, auto-expiry
Offline	node-local ledger, reconcile on reconnect
Latency	< 100 ms in-site, sub-second across channels

Field by

Warehouse systems	3PL & logistics

Schema

```
type Stockunit {
  sku:      string
  node:     Node!  # warehouse or store
  on_hand:  int
  reserved: int  # held for open orders
  version:  int  # deterministic ordering
}
```

A.1 · Interface wireframe
Inventory & Distributed Warehousing Logic

A.2 · Interface wireframe
Hospitality & Unit Management Engines

A.3 · Interface wireframe
Unified Commerce & POS Ledger Sync



B Custom-Engineered · Pillar B
Tailored Commercial Architecture

From-scratch systems, engineered for unique, proprietary operational flows.

B.1
Custom internal operating systems
Replacing problematic spreadsheets.

One system of record in place of a folder of spreadsheets: role-based views, enforced workflows, and a history of every change.

```
record
single source of truth
access
role-based
history
append-only
```

B.2
Private agentic workflows
Operating securely on client infrastructure.

Autonomous logic that runs inside your perimeter, on your hardware or your cloud account. Every action is logged, bounded by explicit permissions, and reversible.

```
runtime
client infrastructure
agents
denied by default
oversight
approvals, full action log
```

B.3
High-availability distributed databases
Data offices to synchronization.

Replicated stores that keep accepting writes when the network does not, then converge deterministically when it returns.

```
writes
local-first
merge
deterministic
single-writer-of-fabric
rules, by design
```

03 System mechanics & technical rigor

Specified like infrastructure, stated so it can be tested.

For CTOs and technical evaluators.

Data ownership
Sovereign
Complete sovereignty. Databases synchronize with client-controlled infrastructure.

Availability
Offline-capable
Operational nodes designed to function without internet disruption.

Latency
< 100 ms
Sub-100ms internal state changes across multi-terminal setups.

Integrations
Deterministic
Clean, deterministic APIs connecting legacy ERPs, payment gateways, and third-party logistics.

Representative event trace — illustrative sequence, not a benchmark

```
# multi-terminal sync, 2-node setup (lan)
t.000 ms   node-01   data_created    sku=…  node=01
t.018 ms   node-01   ledger_posted   28 · 2cr dr
t.031 ms   node-02   replicated      ledger_seq=…
t.407 ms   node-01   state_applied   ✓ (200)
t.412 ms   node-02   state_applied   ✓ (201)

# multi-node, 3-node setup (wan)
t.000 ms   node-01   data_created    sku=…  node=01
t.092 ms   node-02   replicated      ledger_seq=…
t.118 ms   node-03   replicated      ✓
t.…  ms   node-01   state_applied   ✓ (200)
t.…  ms   node-02   replicated      conver-ged ✓ (201)

Target: state changes under 100 ms across terminals, with no dependence on the network.
```

04 Security & compliance

Stated plainly, so it can be checked.

Properties of the systems we build, written as specifications rather than assurances.

We do not claim certifications on this page. If you must meet a specific regulation, we design to it and show you how.

- ✓ **Data location** — Client-controlled infrastructure.
- ✓ **Access** — Role-based. Permissions are explicit and reviewable.
- ✓ **History** — An append-only record of changes and actions.
- ✓ **Agents** — Bounded permissions, approval gates, a full action log, and every action reversible.
- ✓ **Network** — Outbound access denied by default.
- ✓ **Failure** — Offline-capable nodes and deterministic recovery.

05 How we work

Four stages. One owner at each.

Every engagement follows the same order, so you always know what is being decided and what comes next.

01 Understand
We map how work actually moves through your business, systems, hand-offs, and re-entry.

You receive
A written map of current operations, and the baseline for the 30% target.

02 Design
We specify the system: data model, integrations, failure behavior, and what must run offline.

You receive
An architecture document you can hand to your own engineers.

03 Build
We build and test in your environment, in stages, each one usable on its own.

You receive
Working software, tested against the baseline.

04 Operate
We run it with you, monitor it, and hand over ownership when you are ready.

You receive
Runbooks, access, and the option to run it yourselves.

06 Engagement models

Two ways to start.

	Pillar A Turnkey modules	Pillar B Bespoke systems
Best when	Your process matches a standard shape: stock, units, a sales ledger.	Your operations are particular to you and no product fits.
What we do	Configure, connect and deploy a tested module.	Specify and engineer from scratch.
From you	Access to current systems, a named process owner, and your configuration decisions.	Time with the people who run the work, access to systems and data, and a named decision-maker.
You get	A running system, and training.	A system built to your flow, with full documentation.
Start with	The technical dossier.	An architectural walkthrough.

07 Questions from engineers

Plain answers to the usual questions.

Longer answers, with diagrams, are in the technical dossier.

Ask us directly →

Who owns the data?
You do. Databases synchronize with infrastructure you control, and we keep no copy that we depend on.

What happens without internet?

How do we migrate from our current system?

Can it connect to our ERP and payment gateway?

Do agents send our data to outside services?

What if we outgrow a turnkey module?

08 The Sanad manifesto

Software should clarify the mind, not demand it.

Most enterprise software creates its own bureaucracy. Teams spend hours managing the tool that was supposed to save them time. Sanad is built on the inverse thesis: intelligent software works silently in the background so leaders and builders can dedicate their focus to what actually matters.

Support for the people who lead.

If the work is heavier than it should be, we should talk.

Request Technical Dossier → or write to inquiries@sanad.systems

© 2026 Sanad. All rights reserved.
● Accepting new engagements · Remote-first studio

Documentation
Technical documentation
Whitepapers
Operational blueprints
Shared on request, with the technical dossier.

Inquiries
inquiries@sanad.systems
For an architectural walkthrough of any system on this page, write to us directly.